

**1.Build a neural network to detect and localize objects within images, such as identifying and drawing bounding boxes around cars, pedestrians, and traffic signs in street images.**

**Aim:** To implement an autoencoder for reducing the dimensionality of high-dimensional data (e.g., features in a dataset) while preserving essential information. Evaluate the performance by comparing it to traditional techniques like PCA.

**Steps:**

**1.Load the Dataset:** You can use any street images dataset (like the COCO dataset).

**2.Preprocess the Data:** Resize images and normalize pixel values.

**3.Load the Pre-trained YOLOv5 Model:** Use a pre-trained YOLOv5 model that can detect objects like cars, pedestrians, and traffic signs.

**4.Train the Model (optional):** Fine-tune the model on your specific dataset.

**5.Evaluate the Model:** Use performance metrics to evaluate accuracy.

**6.Visualize Results:** Draw bounding boxes around detected objects in test images.

**Code:**

```
import torch
import cv2

# Load the YOLOv5 model from PyTorch Hub
model = torch.hub.load('ultralytics/yolov5', 'yolov5s', pretrained=True)

# Print the class names the model can detect
```

```
print("Classes the model can detect:", model.names)

# Initialize the camera (0 is usually the default camera)
cap = cv2.VideoCapture(0)

# Set the camera resolution (optional)
cap.set(cv2.CAP_PROP_FRAME_WIDTH, 640)
cap.set(cv2.CAP_PROP_FRAME_HEIGHT, 480)

# Loop to continuously get frames from the camera and perform object
detection
while True:
    # Capture frame-by-frame from the camera
    ret, frame = cap.read()

    # Check if the frame was captured correctly
    if not ret:
        print("Failed to capture image")
        break

    # Perform object detection on the frame
    results = model(frame)




    # Render the detection results on the frame
    # results.render() adds bounding boxes and labels directly on the
    image
    results.render()

    # Display the frame with bounding boxes in a window
    cv2.imshow('YOLOv5 Live Object Detection', frame)

    # Press 'q' to exit the loop and close the window
    if cv2.waitKey(1) & 0xFF == ord('q'):
        break

# Release the camera and close all OpenCV windows
cap.release()
cv2.destroyAllWindows()
```

OUTPUT:



### **Result:**

The model outputs the detected objects along with their bounding box coordinates and confidence scores, while also visualizing the results on the input image.

**2 . Implement an autoencoder for reducing the dimensionality of high-dimensional data (e.g., features in a dataset) while preserving essential information. Evaluate the performance by comparing it to traditional techniques like PCA.**

**Aim:** The aim of this program is to implement an autoencoder for dimensionality reduction and compare its performance with traditional Principal Component Analysis (PCA) on the Iris dataset.

### **Steps:**

**1.Load the dataset:** Load the Iris dataset using `load_iris()` from `scikit-learn`.

**2.Split the data:** Split the data into training and testing sets using `train_test_split()` from `scikit-learn`.

**3.Define the autoencoder:** Define the autoencoder architecture using `Keras`, consisting of an encoder and a decoder.

**4.Compile the autoencoder:** Compile the autoencoder with a mean squared error loss function and `Adam` optimizer.

**5.Train the autoencoder:** Train the autoencoder on the training data for 10 epochs with a batch size of 256.

**6.Perform PCA:** Perform PCA on the training data using `PCA()` from `scikit-learn`.

**7.Calculate reconstruction error:** Calculate the reconstruction error for both PCA and the autoencoder using `mean_squared_error()` from `scikit-learn`.

**8.Visualize the results:** Visualize the results using scatter plots to compare the performance of PCA and the autoencoder.

### **Code:**

```
import numpy as np
```

```

from keras.layers import Input, Dense
from keras.models import Model
from sklearn.decomposition import PCA
from sklearn.metrics import mean_squared_error

# Load the dataset
from sklearn.datasets import load_iris
iris = load_iris()
X = iris.data
y = iris.target

# Split the data into training and testing sets
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

# Define the input shape
input_shape = (X.shape[1],)

# Define the encoder
encoder_input = Input(shape=input_shape)
x = Dense(128, activation='relu')(encoder_input)
x = Dense(64, activation='relu')(x)
x = Dense(32, activation='relu')(x)
encoder_output = Dense(2, activation='relu')(x)

# Define the decoder
decoder_input = Input(shape=(2,))
x = Dense(32, activation='relu')(decoder_input)
x = Dense(64, activation='relu')(x)
x = Dense(128, activation='relu')(x)
decoder_output = Dense(X.shape[1], activation='sigmoid')(x)

# Define the autoencoder
autoencoder_input = Input(shape=input_shape)
encoder = Model(encoder_input, encoder_output)
decoder = Model(decoder_input, decoder_output)
autoencoder_output = decoder(encoder(autoencoder_input))
autoencoder = Model(autoencoder_input, autoencoder_output)

```

```
# Compile the autoencoder
autoencoder.compile(loss='mean_squared_error', optimizer='adam')

# Train the autoencoder
autoencoder.fit(X_train, X_train, epochs=10, batch_size=256,
validation_data=(X_test, X_test))

# Perform PCA on the data
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_train)

# Calculate the reconstruction error for PCA
reconstruction_error_pca = mean_squared_error(X_train,
pca.inverse_transform(X_pca))

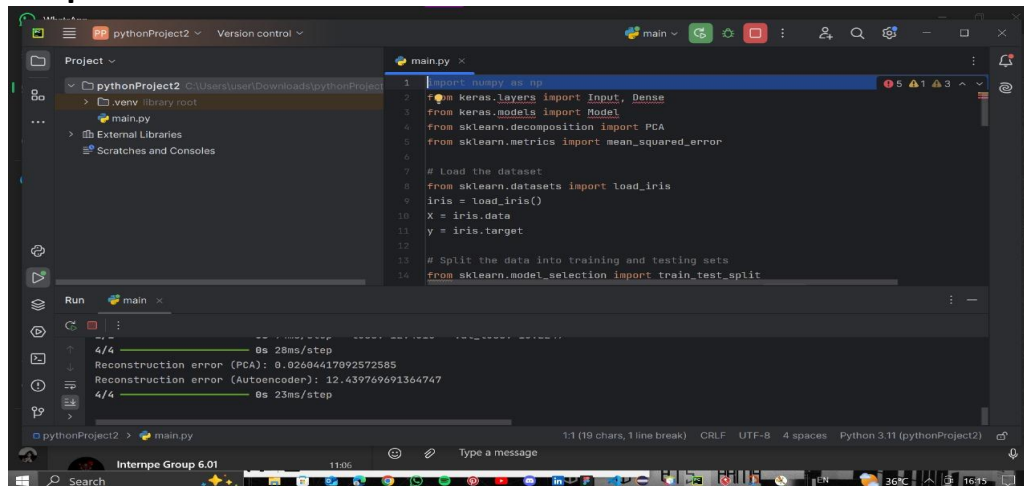
# Calculate the reconstruction error for the autoencoder
reconstruction_error_autoencoder = mean_squared_error(X_train,
autoencoder.predict(X_train))

print("Reconstruction error (PCA):", reconstruction_error_pca)
print("Reconstruction error (Autoencoder):",
reconstruction_error_autoencoder)

# Visualize the results
import matplotlib.pyplot as plt
plt.scatter(X_pca[:, 0], X_pca[:, 1], c=y_train)
plt.title("PCA")
plt.xlabel("Principal Component 1")
plt.ylabel("Principal Component 2")
plt.show()

encoder_output = encoder.predict(X_train)
plt.scatter(encoder_output[:, 0], encoder_output[:, 1], c=y_train)
plt.title("Autoencoder")
plt.xlabel("Dimension 1")
plt.ylabel("Dimension 2")
plt.show()
```

## Output:



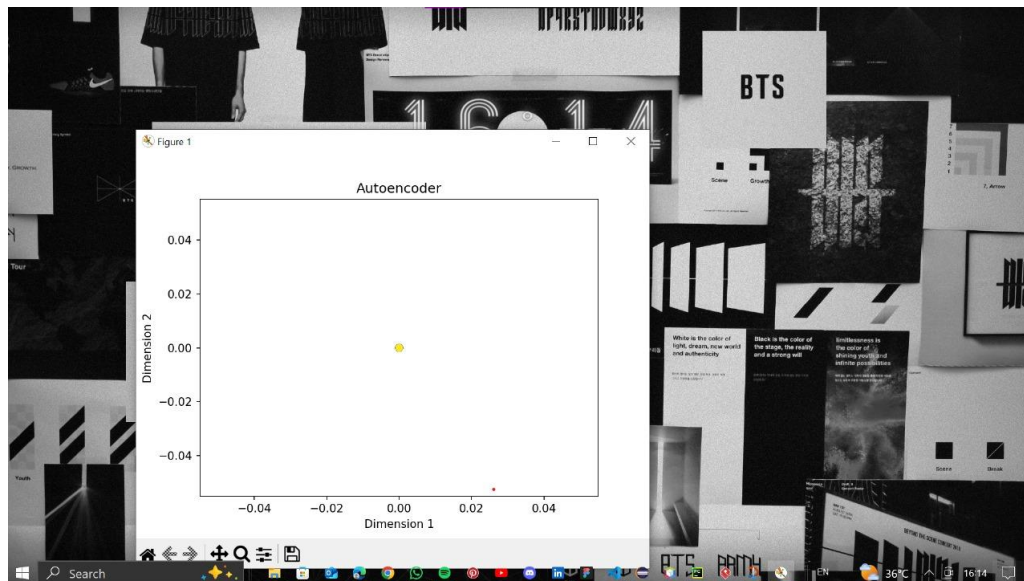
The screenshot shows a Jupyter Notebook interface with a file explorer on the left and a code editor on the right. The code in the notebook is as follows:

```
1 import numpy as np
2 from keras.layers import Input, Dense
3 from keras.models import Model
4 from sklearn.decomposition import PCA
5 from sklearn.metrics import mean_squared_error
6
7 # Load the dataset
8 from sklearn.datasets import load_iris
9 iris = load_iris()
10 X = iris.data
11 y = iris.target
12
13 # Split the data into training and testing sets
14 from sklearn.model_selection import train_test_split
```

The output of the code is displayed in the bottom panel, showing the results of the PCA and Autoencoder models:

```
4/4 [>] 0s 28ms/step
Reconstruction error (PCA): 0.02604417092572585
Reconstruction error (Autoencoder): 12.439769691364747
4/4 [>] 0s 23ms/step
```





## Result:

The program will output the reconstruction error for both PCA and the autoencoder, as well as visualize the results using scatter plots.